

Binary Logistic Regression

Checking Model Performance

# Contents

1. Receiver Operating Characteristic (ROC) Curve
2. Lift Curve
3. Kolmogorov Smirnov statistics
4. Pearson residuals
5. Residual plot
6. Multicollinearity

# Receiver Operating Characteristic Curve

- The Receiver Operating Characteristic (ROC) curve is


A graphical representation of the trade off between the false positive and true positive rates for various cut off values

Y- axis: Sensitivity ( true positive rate)

X-axis: 1-Specificity (false positive rate)

The performance of the classification model can be assessed by area under the ROC curve (C).

# ROC Curve and Area Under ROC Curve



High TPR with low FPR is indicative of a good model. This will result in curve that is closer to the Y-axis and top left corner of the plot. It implies higher Area Under the ROC Curve.

# ROC Curve and Area Under ROC Curve

Interpreting different versions of an ROC curve

Critical Points	Interpretations
TPR = 0 and FPR = 0	Model predicts every instance to be Non-event
TPR = 1 and FPR = 1	Model predicts every instance to be Event
TPR = 1 and FPR = 0	The Perfect Model

- If the model is perfect, AUC = 1
- If the model is guessing randomly, AUC = 0.5
- Thumb rule: Area Under ROC Curve > 0.65 is considered acceptable

# ROC in Python

# Importing bank loan data & Fitting Binary Logistic Regression model

```
import pandas as pd
bankloan=pd.read_csv('BANK LOAN.csv')

bankloan['AGE']=bankloan['AGE'].astype('category')

import statsmodels.formula.api as smf
riskmodel = smf.logit(formula = 'DEFAULTER ~ EMPLOY + ADDRESS +
DEBTINC + CREDDEBT', data = bankloan).fit()

from sklearn.metrics import roc_curve, auc
bankloan=bankloan.assign(pred=riskmodel.predict())
fpr, tpr, thresholds = roc_curve(bankloan['DEFAULTER'],
bankloan['pred'])
```

- ❑ Import **roc\_curve**, **auc** from sklearn.metrics
- ❑ **predict()** function prepares data required for ROC curve.
- ❑ **roc\_curve()** computes Receiver operating characteristic (ROC), it returns "tpr" (True positive rate), "fpr" (False positive rate) and threshold.

# ROC in Python

# AUC & ROC plot

```
ruc_auc = auc(fpr, tpr) ← auc() gives area under curve


import matplotlib.pyplot as plt

plt.figure()
lw = 2
plt.plot(fpr, tpr, color='darkorange', lw=lw, label='ROC curve (area = %0.2f)' % ruc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0]); plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate'); plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic'); plt.legend(loc="lower right")
plt.show()
```

- **plot()** function plots the objects created using roc\_curve. Entire code of plot should be run in single chunk.

# ROC in Python

# Output:



# AUC value

```
print("Area under the ROC curve : %f" % roc_auc)
```

# Output:

```
Area under the ROC curve : 0.855619
```

## Interpretation :


□ Area under the curve is 0.8556 which means model is performing well



# Lift Curve

- The idea is to quantify and compare two scenarios-One is using the model to identify certain cases and second using random selection of cases for specific purpose like marketing campaign.
- Lift is the ratio of results obtained **with and without a model**.
- Although primarily used in marketing analytics, the concept finds applicability in other domains as well, such as risk modeling, supply chain analytics, etc.

# Lift Curve



**Lift Curve:** After contacting X% of customers, Y% of respondents will be identified if statistical model is used.

Ratio  $Y/X$  is plotted

**Baseline:** After contacting X% of customers, X% of respondents will be identified if random method is used.

Ratio  $X/X$  is plotted

# Lift Curve in Python

# Install “scikit-plot” library in Anaconda Prompt and load in Python

```
from sklearn.linear_model import LogisticRegression
import matplotlib.pyplot as plt
import scikitplot as skplt
```

```
X = bankloan[['EMPLOY', 'ADDRESS', 'DEBTINC', 'CREDDEBT']]
y = bankloan[['DEFAULTER']]
log_model = LogisticRegression()
log_model.fit(X,y)
pred_log = log_model.predict_proba(X)

skplt.metrics.plot_lift_curve(y, pred_log)
plt.show()
```

- ❑ **LogisticRegression()**  
fits a Logistic Regression model
- ❑ **predict\_proba()**  
Return probability estimates for the test vector X.
- ❑ **scikitplot()** depends on scikit-learn and




Note : Incase import scikitplot as skplt gives an error, then just run `pip install scikit-plot` in anaconda prompt & then execute import code in python

S

plots for sklearn

# Lift Curve in Python

# Output:



## Interpretation :

- Model is performing better. As more defaulters identified in earlier buckets.

# Kolmogorov-Smirnov Statistic

Kolmogorov-Smirnov (KS) Statistics is one of the most commonly used measures to assess predictive power for marketing or credit risk models

KS is maximum difference between % cumulative Goods and Bads distribution across probability bands

The gains table typically has % cumulative Goods (or Non-Event) and % Cumulative Bads (Or Event) across 10 or 20 probability bands

- **KS is a point estimate**, meaning it is only one value and indicate the probability band where separation between Goods (or Non-Event) and Bads (or Event) is maximum
- Theoretically K-S can range from 0-100. **KS less than 25, may not indicate good model**. Too high value should also be evaluated carefully

# Kolmogorov-Smirnov Statistic

BAND	Count	Percent	Count(bad)	%(bad)	Count(good)	%(good)	cum% bad	cum% good	KS
0.95-1	10	1.4%	9	4.9%	1	0.2%	4.9%	0.2%	4.7%
0.90-0.95	7	1.0%	7	3.8%	0	0.0%	8.7%	0.2%	8.5%
0.85-0.90	7	1.0%	6	3.3%	1	0.2%	12.0%	0.4%	11.6%
0.80-0.85	7	1.0%	5	2.7%	2	0.4%	14.8%	0.8%	14.0%
0.75-0.80	11	1.6%	9	4.9%	2	0.4%	19.7%	1.2%	18.5%
0.70-0.75	17	2.4%	14	7.7%	3	0.6%	27.3%	1.7%	25.6%
0.65-0.70	17	2.4%	12	6.6%	5	1.0%	33.9%	2.7%	31.2%
0.60-0.65	10	1.4%	7	3.8%	3	0.6%	37.7%	3.3%	34.4%
0.55-0.6	24	3.4%	14	7.7%	10	1.9%	45.4%	5.2%	40.1%
0.5-0.55	21	3.0%	9	4.9%	12	2.3%	50.3%	7.5%	42.7%
0.45-0.5	22	3.1%	9	4.9%	13	2.5%	55.2%	10.1%	45.1%
0.40-0.45	31	4.4%	13	7.1%	18	3.5%	62.3%	13.5%	48.8%
0.35-0.4	29	4.1%	11	6.0%	18	3.5%	68.3%	17.0%	51.3%
0.3-0.35	27	3.9%	13	7.1%	14	2.7%	75.4%	19.7%	55.7%
0.25-0.3	40	5.7%	7	3.8%	33	6.4%	79.2%	26.1%	53.1%
0.2-0.25	45	6.4%	12	6.6%	33	6.4%	85.8%	32.5%	53.3%
0.15-0.2	52	7.4%	10	5.5%	42	8.1%	91.3%	40.6%	50.6%
0.10-0.15	66	9.4%	4	2.2%	62	12.0%	93.4%	52.6%	40.8%
0.05-0.1	80	11.4%	8	4.4%	72	13.9%	97.8%	66.5%	31.3%
0-0.05	177	25.3%	4	2.2%	173	33.5%	100.0%	100.0%	0.0%
<b>Total</b>	<b>700</b>	<b>100%</b>	<b>183</b>	<b>100%</b>	<b>517</b>	<b>100%</b>			

# Kolmogorov-Smirnov Statistic in Python

# Combine observed and expected frequencies

```
from scipy.stats import ks_2samp  
ks_2samp(bankloan.loc[bankloan.DEFAULTER==0, 'pred'],  
bankloan.loc[bankloan.DEFAULTER==1, 'pred'])
```

- ❑ **ks\_2samp** computes the kolmogorov-smirnov statistic on 2 samples.
- ❑ It returns KS statistic and two-tailed p-value

# Output:

```
Ks_2sampResult(statistic=0.561552039403452, pvalue=1.909421801103993e-37)
```

## Interpretation :

- ❑ Maximum difference (K-S statistic) is 0.561552.

# Pearson Residuals

- Pearson residual is defined as the standardized difference between the observed and predicted frequency. It measures the relative deviations between the observed and fitted values. :

$$r_j = \frac{(Y_j - M_j p_j)}{\sqrt{M_j p_j (1 - p_j)}}$$

where

$M_j$  : number of observations with  $j$ th covariate pattern

$Y_j$  : Observed value (1 or 0) for  $j$ th covariate pattern

$p_j$  : Predicted probability for  $j^{\text{th}}$  covariate pattern

- Binary Logistic Regression does not require 'Normality' of residuals



# Pearson Residuals in Python

# Obtain residuals

```
bankloan=bankloan.assign(resid=riskmodel.resid_pearson)
```

```
bankloan.head()
```

□ **resid\_pearson()** calculates Pearson residuals.

# Output:

	SN	AGE	EMPLOY	ADDRESS	DEBTINC	CREDDEBT	OTHDEBT	DEFAULTER	pred	resid
0	1	3	17	12	9.3	11.36	5.01	1	0.808346726	0.486921868
1	2	1	10	6	17.3	1.36	4	0	0.198114704	-0.497052463
2	3	2	15	14	5.5	0.86	2.17	0	0.010062815	-0.100822141
3	4	3	15	14	2.9	2.66	0.82	0	0.022159721	-0.150538706
4	5	1	2	0	17.3	1.79	3.06	1	0.781808095	0.528286162


Residuals

# Pearson Residuals Plot in Python

# Residuals Plot

```
import seaborn as sns  
sns.scatterplot('SN', 'resid', data=bankloan); plt.xlabel('SN');  
plt.ylabel('residual')
```

# Output:



Clearly one case has very high residual value.

# Multicollinearity

- Multicollinearity exists if there is a strong linear relationship among the continuous independent variables.
- Do not ignore multicollinearity in Binary Logistic Regression .
- Use variance inflation factors to detect multicollinearity.



**Multicollinearity is explained in MLR module.**

# Quick Recap

In this session, we learnt about **checking model performance** :

ROC	<ul style="list-style-type: none"><li>• Graphical representation of the trade off between the false positive (FPR) and true positive (TPR) rates for various cut off values.</li></ul>
Lift curve	<ul style="list-style-type: none"><li>• Lift Curve Compares model results with baseline without model</li></ul>
K-S statistic	<ul style="list-style-type: none"><li>• KS is the maximum difference between % cumulative Goods (event/<math>Y=1</math>) and cumulative Bads (non events/<math>Y=0</math>) distribution across probability groups.</li></ul>
Residual	<ul style="list-style-type: none"><li>• Pearson's residual is used for binary logistic regression</li></ul>
Multicollinearity	<ul style="list-style-type: none"><li>• Multicollinearity exists if there is a strong linear relationship among the continuous independent variables</li></ul>