

Support Vector Machines in Python

Contents

1. Understanding Association Rules
2. Introduction to Market Basket Analysis
 - i. Uses
 - ii. Definitions and Terminology
3. Rule Evaluation
 - i. Support
 - ii. Confidence
 - iii. Lift
4. Market Basket Analysis in Python

Contents

1. Introduction to Support Vector Machine (SVM)
2. Understanding Hyper Planes
 - i. What is a Hyper Plane
 - ii. Hyper Plane Separation
3. Linear Separators
 - i. Classification Margin
4. Mathematical Approach to Linear SVM
5. Non-Linear SVM
6. About the Kernel Function
7. SVM in Python
 - i. SVM Modeling
 - ii. Confusion Matrix and Area Under ROC Curve

Introduction to Support Vector Machines

- Support Vector Machines (SVM's) are a relatively new learning method generally used for classification problem.
- Although the first paper dates way back to early 1960's it is only in 1992-1995 that this powerful method was universally adopted as a mainstream machine learning paradigm

The basic idea is to find a hyper plane which separates the d -dimensional data perfectly into its classes. However, since training data is often not linearly separable, SVM's introduce the notion of a "Kernel-induced Feature Space" which casts the data into a higher dimensional space where the data is separable.

What is a Hyper Plane

In two dimensions, a hyper plane is defined by the equation:

$$W_1X_1 + W_2X_2 + b = 0$$

This is nothing but equation of line.

The above equation can be easily extended to the p-dimensional setting:

$$W_1X_1 + W_2X_2 + \dots + W_pX_p + b = 0$$


In short,

$$\mathbf{W}^T\mathbf{X} + b = 0$$

In $p > 3$ dimensions, it can be hard to visualize a hyper planes.

Separating a Hyper Plane


- Binary classification can be viewed as the task of separating classes in feature space:




Linear Separators

The objective in SVM is to find optimum separator

Fig. 02: Linear
Separators



Classification Margin




- Distance from case x_i to the separator is

$$r = \frac{\mathbf{w}^T \mathbf{x}_i + b}{\|\mathbf{w}\|}$$

Here $\|\mathbf{w}\|$ is length of a vector given by $\sqrt{\sum(W^2)}$

- Cases closest to the hyper plane are Support Vectors
- Margin ρ of the separator is the distance between support vectors

Maximum Margin Classification



- The objective is now to maximize the margin ρ of the separator
- The focus is on 'Support Vectors'
- Other cases are not considered in the algorithm

Mathematical Approach to Linear SVM

Let training set be separated by a hyper plane with margin ρ . Then for each training observation

$$\begin{aligned} \mathbf{w}^T \mathbf{x}_i + \mathbf{b} &\leq -\rho/2 & \text{if } y_i = -1 \\ \mathbf{w}^T \mathbf{x}_i + \mathbf{b} &\geq \rho/2 & \text{if } y_i = 1 \end{aligned} \quad \Leftrightarrow \quad y_i(\mathbf{w}^T \mathbf{x}_i + \mathbf{b}) \geq \rho/2$$

For every support vector \mathbf{x}_s the above inequality is an equality

After rescaling \mathbf{w} and \mathbf{b} by $\rho/2$ in the equality, we obtain that distance between each \mathbf{x}_s and the hyper plane is

$$r = \frac{y_i(\mathbf{w}^T \mathbf{x}_s + \mathbf{b})}{\|\mathbf{w}\|} = \frac{1}{\|\mathbf{w}\|}$$

Margin can be expressed through (rescaled) \mathbf{w} and \mathbf{b} as:

$$\rho = 2r = \frac{2}{\|\mathbf{w}\|}$$

Mathematical Approach to Linear SVM

Quadratic Optimisation problem is:

Find \mathbf{w} and \mathbf{b} such that

$$\rho = \frac{2}{\|\mathbf{w}\|} \text{ is maximised}$$

and

$$y_i(\mathbf{w}^T \mathbf{x}_i + \mathbf{b}) \geq 1$$

which can be reformulated as:

Find \mathbf{w} and \mathbf{b} such that


$$\phi(\mathbf{w}) = \mathbf{w}^T \mathbf{w} \text{ is minimised}$$

and

$$y_i(\mathbf{w}^T \mathbf{x}_i + \mathbf{b}) \geq 1$$

Non-Linear SVMs – Feature Spaces

General idea: The original feature space can always be mapped to some higher-dimensional feature space where the training set is separable



The "Kernel Trick"

The linear classifier relies on inner product between vectors

$$K(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^T \mathbf{x}_j$$

If every data point is mapped into high-dimensional space via some transformation $\phi: \mathbf{x} \rightarrow \phi(\mathbf{x})$

then the inner product becomes

$$K(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$$

A kernel function is a function that is equivalent to an inner product in some feature space

The "Kernel Trick"

Example:

2-dimensional vector $x = [x_1 \ x_2]$;

Let $K(x_i, x_j) = (1 + x_i^T x_j)^2$

Need to show that $K(x_i, x_j) = \varphi(x_i)^T \varphi(x_j)$:

$$\begin{aligned} K(x_i, x_j) &= (1 + x_i^T x_j)^2 \\ &= 1 + x_{i1}^2 x_{j1}^2 + 2x_{i1}x_{j1}x_{i2}x_{j2} + x_{i2}^2 x_{j2}^2 + 2x_{i1}x_{j1} + 2x_{i2}x_{j2} \\ &= [1 \ x_{i1}^2 \ \sqrt{2}x_{i1}x_{i2} \ x_{i2}^2 \ \sqrt{2}x_{i1} \ \sqrt{2}x_{i2}]^T [1 \\ &\quad x_{j1}^2 \ \sqrt{2}x_{j1}x_{j2} \ x_{j2}^2 \ \sqrt{2}x_{j1} \ \sqrt{2}x_{j2}] \\ &= \varphi(x_i)^T \varphi(x_j) \text{ where } \varphi(x) = [1 \ x_1^2 \sqrt{2}x_1x_2 \ x_2^2 \sqrt{2}x_1 \sqrt{2}x_2] \end{aligned}$$

Thus, a kernel function implicitly maps data to a high-dimensional space (Without the need to compute each $\varphi(x)$ explicitly)

Examples of Kernel Functions

Linear

$$K(x_i, x_j) = x_i^T x_j$$

Mapping ϕ

$x \rightarrow \phi(x)$ where $\phi(x)$ is x itself

Polynomial of power ρ

$$K(x_i, x_j) = (1 + x_i^T x_j)^\rho$$

Gaussian (Radial basis function)

$$K(x_i, x_j) = e^{-\frac{\|x_i - x_j\|^2}{2\sigma^2}}$$

Case Study – Predicting Loan Defaulters

Background

- The bank possesses demographic and transactional data of its loan customers. If the bank has a robust model to predict defaulters it can undertake better resource allocation.

Objective

- To predict whether the customer applying for the loan will be a defaulter

Available Information

- Sample size is 700
- Age group, Years at current address, Years at current employer, Debt to Income Ratio, Credit Card Debts, Other Debts are the independent variables
- **Defaulter** (=1 if defaulter, 0 otherwise) is the dependent variable

Data Snapshot

BANK LOAN							
Independent Variables				Dependent Variable			
SN	AGE	EMPLOY	ADDRESS	DEBTINC	CREDDEBT	OTHDEBT	DEFAULTER
Column	Description	Type	Measurement	Possible Values			
SN	Serial Number	Numeric	-	-			
AGE	Age Groups	Categorical	1(<28 years), 2(28-40 years), 3(>40 years)	3			
EMPLOY	Number of years customer working at current employer	Continuous	-	Positive value			
ADDRESS	Number of years customer staying at current address	Continuous	-	Positive value			
DEBTINC	Debt to Income Ratio	Continuous	-	Positive value			
CREDDEBT	Credit to Debit Ratio	Continuous	-	Positive value			
OTHDEBT	Other Debt	Continuous	-	Positive value			
DEFAULTER	Whether customer defaulted on loan	Binary	1(Default), 0(Non-Defaulter)	2			

SVM in Python

Importing the Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.svm import SVC

from sklearn.metrics import confusion_matrix, f1_score,
precision_score, recall_score, accuracy_score,
roc_curve, roc_auc_score, auc
```

Importing and Readyng the Data

```
bankloan = pd.read_csv("BANK LOAN.csv")

bankloan['AGE'] = pd.Categorical(bankloan['AGE'])

bankloan.info()
bankloan1 = bankloan.drop(['SN','AGE'], axis = 1)
```

❑ **pd.Categorical()**
changes age from an
integer to a factor
variable.

❑ **info()** is used to check if the conversion to category has taken place and if all other variable formats are appropriate, before moving to SVM modeling.

SVM in Python

Output

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 389 entries, 0 to 388  
Data columns (total 8 columns):  
SN            389 non-null int64  
AGE           389 non-null category  
EMPLOY        389 non-null int64  
ADDRESS       389 non-null int64  
DEBTINC       389 non-null float64  
CREDDEBT      389 non-null float64  
OTHDEBT       389 non-null float64  
DEFAULTER     389 non-null int64  
dtypes: category(1), float64(3), int64(4)  
memory usage: 21.9 KB
```

Creating Train and Test Data Sets

```
X = bankloan1.loc[:, bankloan1.columns != 'DEFAULTER']  
y = bankloan1.loc[:, 'DEFAULTER']  
X_train, X_test, y_train, y_test = train_test_split(X, y,  
                                                    test_size=0.30,  
                                                    random_state = 999)
```

- ☐ **train_test_split()** from sklearn.model_selection is used to split dataset into random train and test sets.
- ☐ **test_size** represents the proportion of dataset to be included in the test set.
- ☐ **random_state** sets the seed for the random number generator.

SVM in Python

Model fitting

```
svclassifier = SVC(kernel='linear', probability=True)  
svclassifier.fit(X_train, y_train)
```

- ❑ **svc()** trains a support vector machine.
- ❑ **kernel=** specifies the kernel type to be used in the algorithm ('linear', 'poly', 'rbf', 'sigmoid', 'precomputed').

Output

```
SVC(kernel='linear', probability=True)
```

Predicted Probabilities

```
predprob_test = svclassifier.predict_proba(X_test)
```

- ❑ **predict_proba()** returns predicted probabilities for the test data.

Predictions Based on SVM

Custom Cutoff Value for Prediction Labels

```
cutoff = 0.3  
pred_test = np.where(predprob_test[:,1] > cutoff, 1, 0)  
pred_test
```

□ **np.where** is applied to the probabilities of true event to consider custom cutoff value. The output is an array of binary labels.

Output

```
array([0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 1, 0, 1,  
       1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0,  
       0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1,  
       0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0,  
       0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0,  
       0, 0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0,  
       0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,  
       1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 0, 0, 0,  
       1, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1,  
       0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1])
```

Confusion Matrix and Area Under ROC Curve

Confusion Matrix

```
confusion_matrix(y_test, pred_test, labels=[0, 1])
```

```
array([[118,  36],  
       [ 13,  43]])
```

```
accuracy_score(y_test, pred_test)  
0.7666666666666667
```

```
precision_score(y_test, pred_test)  
0.5443037974683544
```

```
recall_score(y_test, pred_test)  
0.7678571428571429
```

- ❑ **accuracy_score()** = number of correct predictions out of total predictions
- ❑ **precision_score()** = true positives / (true positives + false positives)
- ❑ **recall_score()** also known as 'Sensitivity' = true positives / (true positives + false negatives)

Area Under ROC Curve

```
auc = roc_auc_score(y_test, predprob_test[:,1])  
print('AUC: %.3f' % auc)  
AUC: 0.847
```



Note : Output will be slightly different as observations are randomly assigned to train-test data.

ROC Curve and Area Under ROC Curve

ROC Curve

```
fpr, tpr, thresholds = roc_curve(y_test, predprob_test[:,1])


#Compute AUC using 'auc' function
roc_auc = auc(fpr, tpr)

#Plot the curve for model

plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area =
%0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
plt.legend(loc="lower right")
plt.show()
```

ROC Curve and Area Under ROC Curve

Output:



Quick Recap

In this session, we learnt about **Support Vector Machines**:

Support Vector Machines

- SVMs find a hyper plane which separates the d-dimensional data perfectly into its classes
- Since training data is often not linearly separable, SVM's introduce the notion of a "Kernel-induced Feature Space" which casts the data into a higher dimensional space where the data is separable

SVM in Python

- Library "**sklearn.svm**" has **SVC()** that trains a support vector machine
- The function takes arguments to specify whether **SVC()** is to be used for classification or regression; if probabilities are to be returned and which kernel to use for training and predicting